



Open-weight IAA evaluation — STEMO ambiguous-referent benchmark (1,056 questions)


Overall IAA score




IAA score vs. ambiguity level



Strict-K accuracy by ambiguity type



Behavioral metric breakdown



—●— InternVL3.5-8B —●— Qwen3.5-27B —●— Qwen3.6-27B —●— Qwen3-VL-32B

InternVL3.5-38B omitted: its -HF checkpoint collapses to single-commit answers (clarifies 0.3% vs the 8B's 14%); verified independent of tensor-parallel size, so treated as an invalid run pending re-evaluation.