




0:12 attempt 1

Question, $K = 3$ readings



0:58 attempt 2

Model's first response



3:26 attempt 3

Per-reading credit



Does the ball go
into the door?

gold readings and answers

1 first attempt no

2 second attempt intended yes

3 third attempt yes

Enumerate (singly or grouped)

LLM judge assigns each reading

"Every attempt succeeds^{2,3} except the first.¹"

1 no

2 yes

3 yes

3/3

Clarify

scope-anchored "Which attempt do you mean?"

vague "Could you be more specific?"

interlocutor: "The second attempt." model: "Yes."

1 -

2 yes

3 -

1.0

0.5

only the intended reading is answered

Commit silently

"Yes."

1 ?

2 ?

3 ?

0

ReQueST score = mean credit per question

credit badges show the per-response score